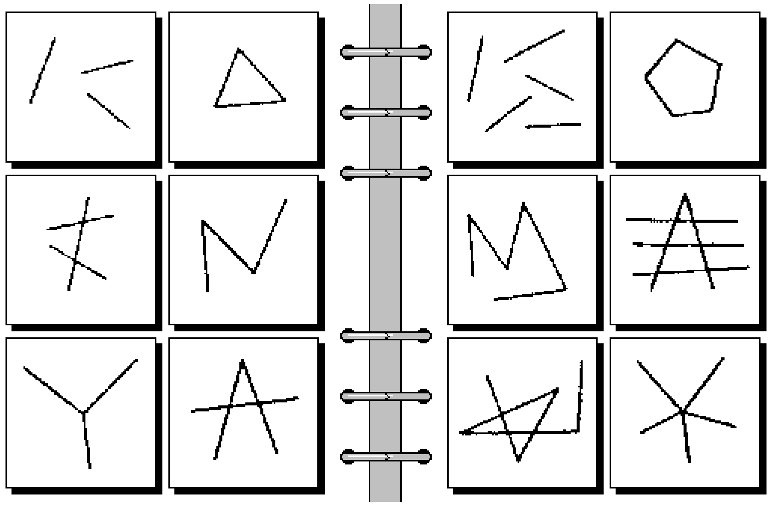

Figure 12: **Synthetic BP #85.** This is the only BP instance for which the voting-based evaluation of MLLM solutions differed from our manual evaluation. Left: Three parts. Right: Five parts. GPT-4o answer: *Left: All images are composed of exactly three lines. Right: All images are composed of more than three lines.* Voting marked it as incorrect, whereas in manual evaluation it was marked as correct. The first difference lies in the meaning of the words *lines* and *parts*, which, in this visual context, seems identical. The second difference stems from the number of the parts on the right side of the problem. The answer seems to be correct, as obviously, five is more than three. However, one could argue that the answer is incomplete, as each of the squares on the right side clearly depicts exactly five parts.

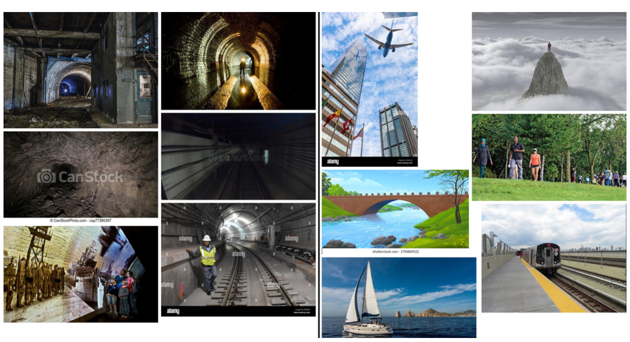

(a) Left: Underground tunnels beneath the city. Right: NOT Underground tunnels beneath the city. GPT-4o solution: Left: *All images depict scenes that are primarily indoors or underground.* Right: *All images depict scenes that are primarily outdoors*. We evaluated it as correct, while the voting marked it as incorrect. The difference arises from how the left concept is perceived. Although the images on the left depict an underground setting, they do not appear to represent an indoor scene. Nevertheless, one of the statements is still true.

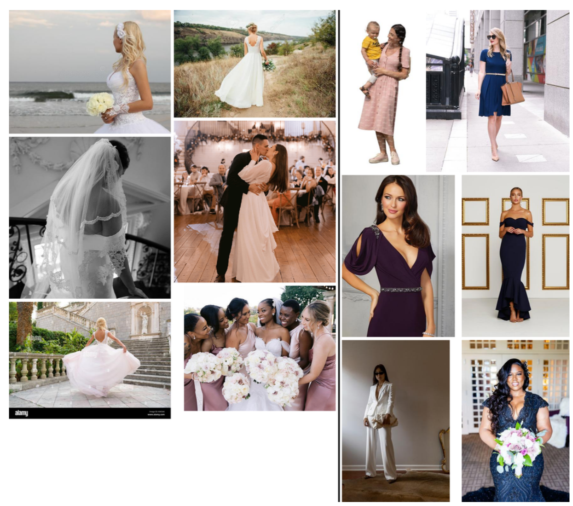

(b) Left: A woman wearing a white wedding dress. Right: NOT A woman wearing a white wedding dress. GPT-4o solution: Left: *All images feature women in white wedding dresses or wedding-related scenes.* Right: *All images feature women in non-wedding attire, wearing dresses or suits of various colors other than white*. We evaluated it as incorrect, while the voting marked it as correct. The difference stems from a small detail in the right concept. One of the images on the right depicts a woman in a white suit, which conflicts with the model's answer.

Figure 13: **The only two Bongard-OpenWorld problems (out of the selected 20) for which the voting evaluation differed from our manual evaluation.** The correctness of GPT-4o's solutions to these problems is disputable.

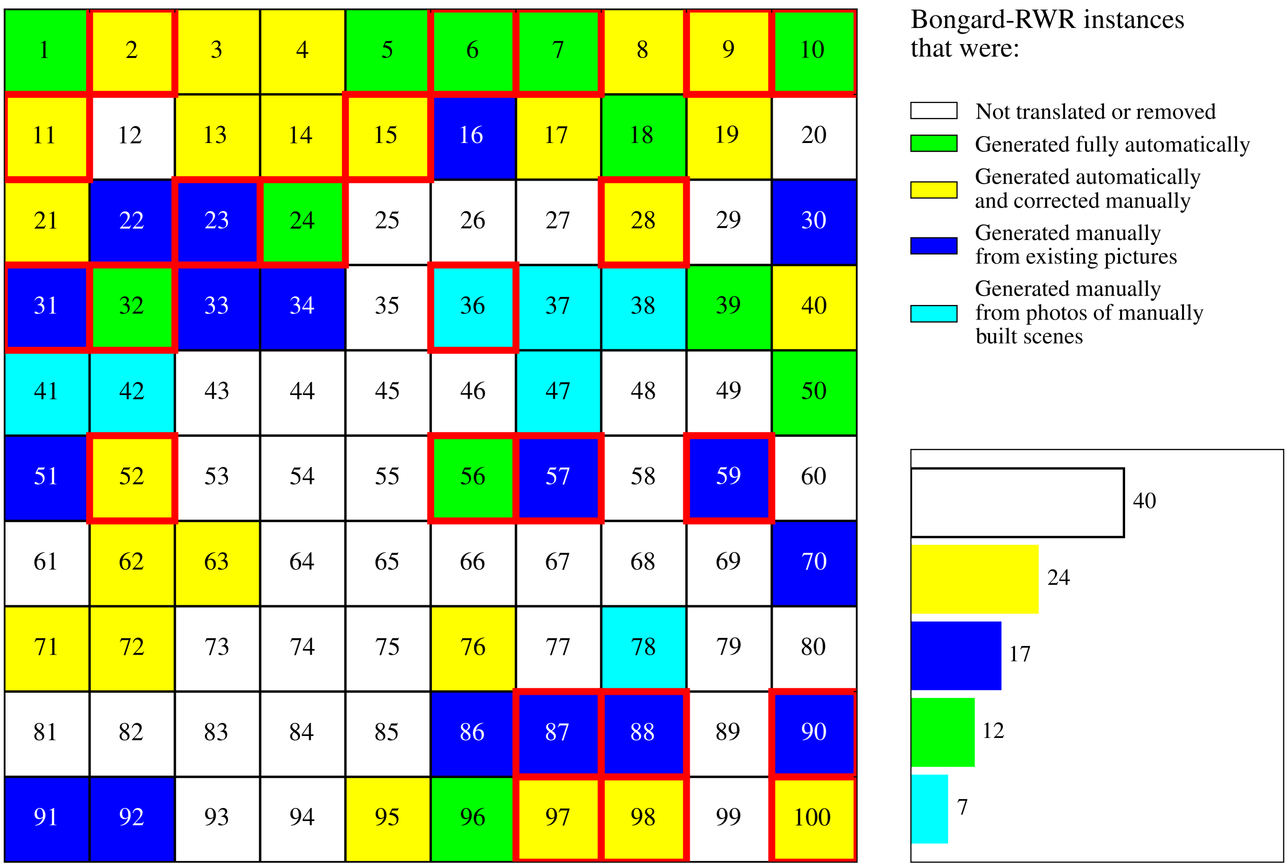

Figure 14: **The structure of the Bongard-RWR dataset.** Each color denotes the genesis of the translation of the respective problem. Problems not translated and those rejected after translation are combined into one category. Red color outlines denote problems which were solved by any model using any strategy. There is no visible correlation between the set of solved problem and the methods used for their generation.

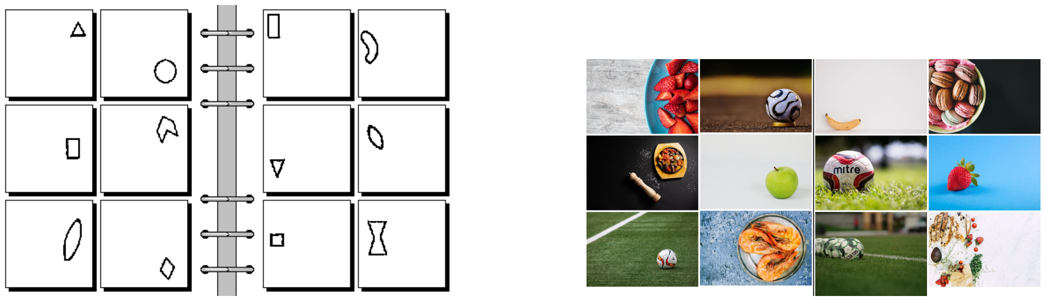

(a) Synthetic BP #8 with its automatically translated Bongard-RWR version. Left: Figures on the right side. Right: Figures on the left side.

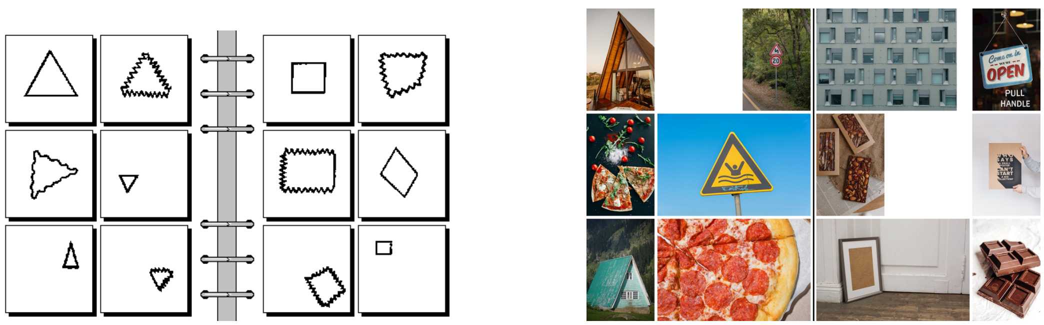

(b) Synthetic BP #10 with its automatically translated Bongard-RWR version. Left: Triangles. Right: Quadrangles.

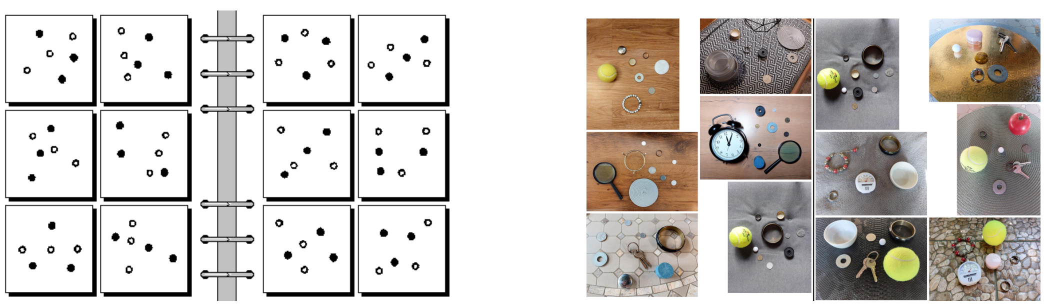

(c) Synthetic BP #41 with it manually constructed Bongard-RWR version. Left: Outline circles on one straight line. Right: Outline circles not on one straight line.

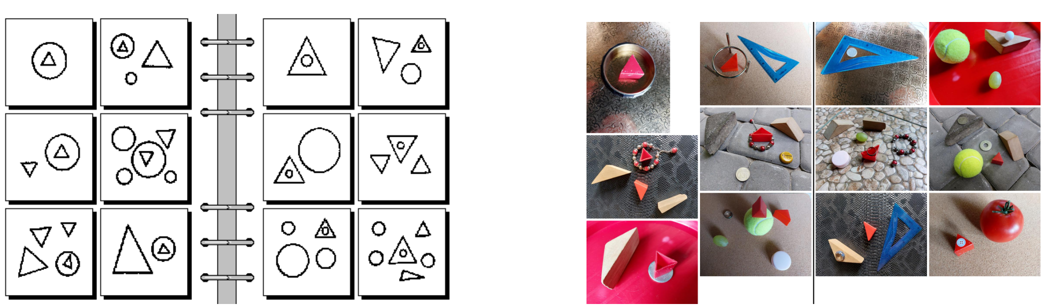

(d) Synthetic BP #47 with its manually constructed Bongard-RWR version. Left: Triangle on top of the circle. Right: Circle on top of the triangle.

Figure 15: **Additional examples of Bongard-RWR instances.**